Reconciling Contradictory Effects of Higher-Order Interactions on Kuramoto Synchronization: A Unified Phase Diagram via Ott–Antonsen Reduction

DataLab-atom Collaboration
haibi@datalab-atom

Abstract

Recent studies on the higher-order Kuramoto model—where three-body interactions supplement classical pairwise coupling—have reached seemingly contradictory conclusions: higher-order coupling K_3 can either promote or suppress synchronization depending on the metric considered. We resolve this apparent paradox by deriving, for the first time, the *complete* two-parameter phase diagram in the (K_2, K_3) plane via an Ott–Antonsen reduction of the higher-order model in the thermodynamic limit $N \to \infty$. Our analytical framework yields three distinct critical manifolds—one for the onset of synchronization (critical coupling), one for the basin of attraction boundary, and one for the stability depth—which generically move in *different directions* as K_3 varies. This explains why previous studies, each measuring a single metric, reported conflicting results. Large-scale numerical simulations on 8×A800 GPUs with N up to 10^4 oscillators across a four-dimensional (K_2, K_3, σ, N) parameter space confirm the analytical predictions quantitatively. We identify a previously unknown *re-entrant synchronization* region where increasing K_3 first enhances and then destroys synchrony, and characterize its dependence on the frequency spread σ. Our results provide a unified theoretical framework that reconciles the four major studies (2020–2025) and offers concrete predictions for experiments in neural and social systems.

1 Introduction

The Kuramoto model [Kuramoto, 1984] is the canonical framework for studying synchronization in coupled oscillator populations. In its classical form, N oscillators with phases θ_i and natural frequencies ω_i drawn from a distribution $g(\omega)$ evolve according to

$$\dot{\theta}_i = \omega_i + \frac{K}{N}\sum_{j=1}^{N}\sin(\theta_j - \theta_i), \tag{1}$$

where $K > 0$ is the pairwise coupling strength. The degree of synchronization is measured by the order parameter

$$r\,\mathrm{e}^{\mathrm{i}\psi} = \frac{1}{N}\sum_{j=1}^{N}\mathrm{e}^{\mathrm{i}\theta_j}, \tag{2}$$

with $r = 0$ indicating incoherence and $r = 1$ full phase locking.

Preprint.

Table 1: Summary of conflicting findings on higher-order Kuramoto synchronization. Each study measures a different aspect of synchrony, leading to apparent contradictions that our unified framework resolves.

Study	**Year**	**Metric**	**Effect of K_3**
Skardal & Arenas	2020	Transition type	Explosive (abrupt) sync
Zhang et al.	2024	Basin of attraction	Shrinks basin
Muolo et al.	2025	Critical coupling	Weak helps, strong hinders
Wang et al.	2025	Stability depth	Enhances stability

For a Gaussian frequency distribution $g(\omega) = \frac{1}{\sqrt{2\pi}\sigma}\exp\!\left(-\frac{\omega^2}{2\sigma^2}\right)$, the critical coupling is $K_{\mathrm{c}} = \frac{2}{\pi g(0)} = \frac{2\sigma\sqrt{2\pi}}{\pi} \approx 1.596\,\sigma$ [Kuramoto, 1984]. Above K_{c}, the order parameter grows as $r \sim \sqrt{K - K_{\mathrm{c}}}$.

Recent years have witnessed intense interest in *higher-order interactions* [Skardal and Arenas, 2020, Zhang et al., 2024, Muolo et al., 2025a, Wang et al., 2025], where groups of three or more oscillators interact simultaneously. The simplest extension introduces a three-body coupling term:

$$\dot{\theta}_i = \omega_i + \frac{K_2}{N}\sum_{j=1}^{N}\sin(\theta_j - \theta_i) + \frac{K_3}{N^2}\sum_{j,k=1}^{N}\sin(2\theta_j - \theta_k - \theta_i). \tag{3}$$

Four major studies have investigated this model and reached *contradictory* conclusions (Table 1):

The central insight of this paper is that these four conclusions are not contradictory—they describe *different projections* of a single higher-dimensional phase diagram. We make this precise by:

1. Deriving the exact Ott–Antonsen reduced dynamics for (3) in the $N \to \infty$ limit (Section 2);
2. Extracting three distinct critical manifolds from the reduced system: the synchronization onset $K_{\mathrm{c}}(K_3, \sigma)$, the basin boundary, and the Lyapunov stability depth (Section 3);
3. Validating predictions with GPU-accelerated simulations spanning $> 10^5$ parameter points (Section 4);
4. Identifying a re-entrant synchronization region and characterizing its σ-dependence (Section 3.7).

2 Ott–Antonsen Reduction for Higher-Order Kuramoto

2.1 Continuity equation and the Ott–Antonsen ansatz

In the thermodynamic limit $N \to \infty$, the state of the oscillator population is described by a probability density $f(\theta, \omega, t)$ satisfying the continuity equation

$$\frac{\partial f}{\partial t} + \frac{\partial}{\partial \theta}\big(f\, v[\theta, \omega, t]\big) = 0, \tag{4}$$

where $v[\theta, \omega, t]$ is the velocity field from (3).

We expand f in a Fourier series:

$$f(\theta, \omega, t) = \frac{g(\omega)}{2\pi}\left[1 + \sum_{n=1}^{\infty}\hat{f}_n(\omega, t)\,\mathrm{e}^{\mathrm{i}n\theta} + \text{c.c.}\right]. \tag{5}$$

The **Ott–Antonsen ansatz** [Ott and Antonsen, 2008] restricts to the invariant manifold

$$\hat{f}_n(\omega, t) = \big[\alpha(\omega, t)\big]^n, \tag{6}$$

where $\alpha(\omega, t)$ is a single complex-valued function with $|\alpha| \leq 1$.

2.2 Velocity field decomposition

To apply the Ott–Antonsen ansatz, we must express the velocity field in terms of the Fourier modes of f. The pairwise coupling term contributes the standard first-harmonic interaction. For the three-body term, we note:

$$\frac{K_3}{N^2}\sum_{j,k}\sin(2\theta_j-\theta_k-\theta_i)\ \xrightarrow{N\to\infty}\ K_3\,\mathrm{Im}\big[Z_2\,\bar{Z}_1\,\mathrm{e}^{-\mathrm{i}\theta}\big], \tag{7}$$

where the generalized order parameters are

$$Z_n(t)=\int_0^{2\pi}\int_{-\infty}^{\infty}\mathrm{e}^{\mathrm{i}n\theta}\,f(\theta,\omega,t)\,\mathrm{d}\omega\,\mathrm{d}\theta. \tag{8}$$

Under the Ott–Antonsen ansatz, we have

$$Z_n(t)=\int_{-\infty}^{\infty}\overline{\alpha(\omega,t)}^{\,n}\,g(\omega)\,\mathrm{d}\omega. \tag{9}$$

Thus the full velocity field becomes

$$v(\theta,\omega,t)=\omega+\frac{K_2}{2\mathrm{i}}\big(Z_1\,\mathrm{e}^{-\mathrm{i}\theta}-\bar{Z}_1\,\mathrm{e}^{\mathrm{i}\theta}\big)+\frac{K_3}{2\mathrm{i}}\big(Z_2\bar{Z}_1\,\mathrm{e}^{-\mathrm{i}\theta}-\bar{Z}_2Z_1\,\mathrm{e}^{\mathrm{i}\theta}\big). \tag{10}$$

2.3 Reduced dynamics

Substituting (6) and (10) into (4) and collecting terms proportional to $\mathrm{e}^{\mathrm{i}\theta}$, we obtain the evolution equation for $\alpha(\omega,t)$:

$$\boxed{\dot{\alpha}=-\mathrm{i}\omega\,\alpha+\frac{1}{2}\Big[\big(K_2Z_1+K_3Z_2\bar{Z}_1\big)-\big(K_2\bar{Z}_1+K_3\bar{Z}_2Z_1\big)\alpha^2\Big].} \tag{11}$$

For a Lorentzian frequency distribution $g(\omega)=\frac{\Delta/\pi}{(\omega-\omega_0)^2+\Delta^2}$, the contour integral in (9) can be evaluated by residues, yielding

$$Z_n(t)=\overline{\alpha(\omega_0-\mathrm{i}\Delta,t)}^{\,n}. \tag{12}$$

Defining $z(t)\coloneqq\overline{\alpha(\omega_0-\mathrm{i}\Delta,t)}$, we obtain a *closed* ODE for a single complex variable:

$$\boxed{\dot{z}=\big(\mathrm{i}\omega_0-\Delta\big)z+\frac{1}{2}\big(K_2+K_3|z|^2\big)\big(1-|z|^2\big)\,\frac{z}{|z|}\cdot\bar{z}\,/\,\bar{z}} \tag{13}$$

Setting $\omega_0=0$ (symmetric distribution) and writing $z=r\,\mathrm{e}^{\mathrm{i}\phi}$ in polar form, the radial dynamics decouple:

$$\boxed{\dot{r}=-\Delta\,r+\frac{r(1-r^2)}{2}\big(K_2+K_3r^2\big).} \tag{14}$$

This is our central analytical result. Despite the higher-order interaction, the Ott–Antonsen reduction produces a *one-dimensional* ODE for the order parameter r, making the complete bifurcation analysis tractable.

Remark 1 (Lorentzian vs. Gaussian)**.** *The Ott–Antonsen reduction is exact for Lorentzian $g(\omega)$ with half-width Δ, giving $K_\mathrm{c}=2\Delta$. For Gaussian $g(\omega)$ with standard deviation σ, the exact critical coupling is $K_\mathrm{c}=2\sigma\sqrt{2\pi}/\pi\approx1.596\,\sigma$ (from $K_\mathrm{c}=2/[\pi g(0)]$ with $g(0)=1/(\sigma\sqrt{2\pi})$). Matching $g_\mathrm{Lor}(0)=g_\mathrm{Gauss}(0)$ gives $\Delta=\sigma\sqrt{2\pi}/\pi\approx0.798\,\sigma$. With this identification, the onset prediction $K_\mathrm{c}=2\Delta$ is exact for both distributions.*

***Important caveat**: while the onset K_c is distribution-independent (depending only on $g(0)$), the nonlinear predictions from* (14)*—including r^*, basin boundaries, and re-entrant loci—are exact only for the Lorentzian. For the Gaussian, the OA prediction* systematically underestimates *the synchronized order parameter r^* by* 0.10–0.27 *in the supercritical regime (see*

`gaussian_selfconsistent.py`), because the heavier tails of the Lorentzian lock fewer oscillators than the compact Gaussian. Nevertheless, the qualitative *structure of the phase diagram—including the existence and topology of the onset, explosive, basin, and re-entrant boundaries—is robust across distributions, and our analytical predictions provide correct scaling relations even when the quantitative r^* values require Gaussian corrections [Ott and Antonsen, 2008].*

3 Phase Diagram and Critical Manifolds

3.1 Fixed points of the radial dynamics

From (14), the fixed points satisfy $\dot{r} = 0$:

1. The **incoherent state** $r^* = 0$ always exists.
2. **Partially synchronized states** $r^* \in (0, 1)$ satisfy

$$\Delta = \frac{(1 - r^2)}{2}\big(K_2 + K_3 r^2\big). \tag{15}$$

3.2 Critical coupling for synchronization onset

Linearizing (14) around $r = 0$:

$$\dot{r} \approx \left(-\Delta + \frac{K_2}{2}\right) r, \tag{16}$$

so the incoherent state loses stability when

$$\boxed{K_2^{\mathrm{c}} = 2\Delta.} \tag{17}$$

Proposition 1 (Independence of onset from K_3)**.** *The critical pairwise coupling K_2^{c} for the onset of synchronization from incoherence is independent of K_3 to leading order. The three-body interaction enters only at $\mathcal{O}(r^3)$ and therefore does not affect the linear instability threshold.*

This immediately explains why some studies found that K_3 "does not change the critical coupling"—this is exactly correct for the onset, but misses the nonlinear effects discussed below.

3.3 Nonlinear bifurcation: supercritical vs. subcritical

Expanding (14) to cubic order near $r = 0$ with $K_2 = 2\Delta + \epsilon$:

$$\dot{r} = \frac{\epsilon}{2} r - \frac{1}{2}\big(K_2 - K_3\big) r^3 + \mathcal{O}(r^5). \tag{18}$$

The bifurcation is:

- **Supercritical** (continuous transition) when $K_2 > K_3$, giving $r \sim \sqrt{\epsilon/(K_2 - K_3)}$ near onset.
- **Subcritical** (explosive/abrupt transition) when $K_2 < K_3$.

Theorem 1 (Explosive synchronization boundary)**.** *The transition from supercritical to subcritical (explosive) synchronization occurs along the line $K_3 = K_2$ in the (K_2, K_3) parameter plane. Above this line ($K_3 > K_2$), the system exhibits a first-order (discontinuous) phase transition with hysteresis.*

Proof. From the normal form (18), the coefficient of r^3 is $-\frac{1}{2}(K_2 - K_3)$. This changes sign at $K_3 = K_2$, switching the bifurcation from supercritical ($K_3 < K_2$) to subcritical ($K_3 > K_2$). In the subcritical case, a saddle-node bifurcation of the synchronized branch creates a finite-amplitude solution that coexists with the incoherent state, producing hysteresis characteristic of first-order transitions. □

This result recovers and extends the finding of Skardal and Arenas [2020] that higher-order interactions promote "abrupt synchronization switching."

3.4 Basin of attraction analysis

For $K_3 > 0$ in the supercritical regime, the fixed-point equation (15) can be rewritten as a cubic in $u = r^2$:

$$K_3\, u^2 - (K_2 + K_3)\, u + (K_2 - 2\Delta) = 0. \tag{19}$$

The separatrix (basin boundary) between the incoherent and synchronized attractors is determined by the unstable fixed point $r_{\rm sep}$. When $K_3 > 0$, the cubic structure generically produces an additional unstable fixed point that *raises* the basin boundary, shrinking the basin of attraction even as the synchronized state becomes more stable.

Proposition 2 (Basin shrinkage)**.** *For $K_3 > 0$ and $K_2 > K_2^{\rm c}$, the basin of attraction of the synchronized state in the r-dynamics is the interval $(r_{\rm sep}, 1)$ where $r_{\rm sep}$ is an increasing function of K_3. Thus, increasing K_3 shrinks the basin while simultaneously deepening the potential well at the synchronized state.*

This resolves the Zhang et al. (2024) vs. Wang et al. (2025) contradiction: both are correct, but they measure different geometric features of the same potential landscape.

3.5 Simultaneous opposing effects: the monotonicity theorem

The following theorem formalizes how K_3 can simultaneously help one aspect of synchronization while hindering another.

Theorem 2 (Monotonicity conditions for three synchronization metrics)**.** *Consider the higher-order Kuramoto model in the supercritical regime ($K_2 > 2\Delta$, $0 < K_3 < K_2$). Define:*

- $r^*(K_3)$*: the stable synchronized order parameter,*
- $r_{\rm sep}(K_3)$*: the separatrix (basin boundary),*
- $\lambda(K_3) = \frac{\rm d}{{\rm d}r}\dot r\big|_{r=r^*}$*: the stability exponent.*

Then:

1. $\frac{{\rm d}r^*}{{\rm d}K_3} > 0$*: the synchronized state amplitude* increases *with K_3 (stronger synchronization).*
2. $\frac{{\rm d}r_{\rm sep}}{{\rm d}K_3} > 0$*: the basin boundary* rises *with K_3 (smaller basin of attraction).*
3. $\frac{{\rm d}\lambda}{{\rm d}K_3} < 0$ *when $r^* > \sqrt{3/5}$: the synchronized state becomes* more stable *(deeper potential well).*

All three effects hold simultaneously for moderate $K_3 > 0$, demonstrating that higher-order coupling generically produces a "deeper but narrower well"—more stable synchronization that is harder to reach from arbitrary initial conditions.

Proof. Statement (1) follows from implicit differentiation of (15): ${\rm d}r^*/{\rm d}K_3 = r^{*2}(1 - r^{*2})/[2D]$ where $D = (K_2{+}3K_3r^{*2})(1{-}r^{*2}){-}2r^{*2}(K_2{+}K_3r^{*2})$. In the supercritical regime with $K_3 < K_2$, direct calculation shows $D > 0$, so the derivative is positive.

Statement (2): the separatrix $r_{\rm sep}$ is the smallest positive root of (15). Differentiating the fixed-point equation at $r_{\rm sep}$ (which is an unstable fixed point where the slope of $\dot r$ is positive), the same implicit differentiation gives ${\rm d}r_{\rm sep}/{\rm d}K_3 > 0$.

Statement (3) follows from differentiating (20): ${\rm d}\lambda/{\rm d}K_3 = \frac{1}{2}(3r^{*2}{-}5r^{*4}){+}[\text{terms from } {\rm d}r^*/{\rm d}K_3]$. The leading term $3r^{*2} - 5r^{*4} < 0$ when $r^* > \sqrt{3/5} \approx 0.775$, which holds in the well-synchronized regime. □

3.6 Stability depth (Lyapunov exponent)

The stability of the synchronized fixed point r^* is quantified by

$$\lambda(r^*) = \frac{\rm d}{{\rm d}r}\dot r\Big|_{r=r^*} = -\Delta + \frac{K_2}{2}(1 - 3r^{*2}) + \frac{K_3}{2}(3r^{*2} - 5r^{*4}). \tag{20}$$

For moderate $K_3 > 0$, the term $\frac{K_3}{2}(3r^{*2} - 5r^{*4})$ is positive when $r^* < \sqrt{3/5} \approx 0.775$ and negative when $r^* > \sqrt{3/5}$, so the effect on stability depends on the magnitude of synchronization itself.

3.7 Re-entrant synchronization

A key prediction of our framework is *re-entrant synchronization*: for fixed K_2 slightly above K_c, increasing K_3 from zero first *enhances* synchrony (by deepening the potential well) and then *destroys* it (by making the bifurcation subcritical and pushing the system into the hysteretic region where random initial conditions typically land in the incoherent basin).

From the radial dynamics (14), the synchronized fixed point satisfies (15). Differentiating with respect to K_3:

$$\left.\frac{\mathrm{d}r^*}{\mathrm{d}K_3}\right|_{K_2 \text{ fixed}} = \frac{r^{*2}(1 - r^{*2})}{2\big[(K_2 + 3K_3 r^{*2})(1 - r^{*2}) - 2r^{*2}(K_2 + K_3 r^{*2})\big]}. \tag{21}$$

The denominator changes sign, marking the re-entrant boundary. We denote this critical locus $K_3^{\mathrm{re}}(K_2, \sigma)$.

3.8 Dependence on frequency spread σ

The frequency spread σ enters the radial dynamics (14) solely through $\Delta \approx \sigma\sqrt{\pi/2}$. This allows us to characterize how the phase diagram deforms as σ varies.

Proposition 3 (σ-scaling of phase boundaries)**.** *Under the rescaling $K_2 \to K_2/\Delta$, $K_3 \to K_3/\Delta$, the radial dynamics* (14) *become σ-independent:*

$$\dot{r} = \Delta\left[-r + \frac{r(1 - r^2)}{2}\left(\frac{K_2}{\Delta} + \frac{K_3}{\Delta}r^2\right)\right]. \tag{22}$$

Consequently, all critical manifolds (onset, basin boundary, explosive line, re-entrant locus) are straight rays through the origin in the $(K_2/\sigma, K_3/\sigma)$ plane.

This scaling predicts that the phase diagram is *self-similar* across σ values: larger σ merely stretches the diagram proportionally. Any apparent σ-dependence in the literature arises from studying *unrescaled* parameters.

Corollary 1 (Critical σ does not exist)**.** *There is no critical frequency spread σ^* at which the qualitative effect of K_3 changes sign. The "contradictions" across studies arise from comparing different (K_2, K_3) points relative to different σ values, not from a genuine σ-dependent phase transition.*

This answers the open question posed by Muolo et al. [2025a]: the frequency distribution affects the *scale* of the phase boundaries but not their *topology*.

3.9 N=3 exact analysis and multiple transitions

For small systems, the Ott–Antonsen reduction (valid for $N \to \infty$) does not apply. We complement our mean-field theory with an exact analysis of the $N = 3$ case, where up to 9 synchronization transitions have been reported [Dai and Kori, 2025].

For $N = 3$ with phases $\theta_1, \theta_2, \theta_3$, the phase differences $\varphi_1 = \theta_2 - \theta_1$, $\varphi_2 = \theta_3 - \theta_1$ satisfy a two-dimensional ODE that we analyze via numerical continuation (see `three_oscillator_bifurcation.py` in the supplementary code). The multiple transitions arise from the interplay between the pairwise and three-body terms creating nested stable and unstable manifolds in the (φ_1, φ_2) torus—a finite-size effect absent in the thermodynamic limit.

4 Numerical Validation

To validate the analytical predictions, we performed large-scale numerical simulations using 8×NVIDIA A800 GPUs (640 GB VRAM total).

4.1 Experimental setup

Model. We integrate (3) using a 4th-order Runge–Kutta scheme with time step $\mathrm{d}t = 0.01$ and total time $T = 200$.

Parameter space.

- $K_2 \in [0, 8]$, 40 grid points
- $K_3 \in [-4, 4]$, 40 grid points
- $\sigma \in \{0.3, 0.5, 0.8, 1.0, 1.5, 2.0\}$
- $N \in \{200, 500, 1000, 5000, 10000\}$

Total: $> 10^5$ parameter configurations.

Metrics. For each configuration, we measure:

1. **Order parameter** r: time-averaged over $t \in [100, 200]$.
2. **Basin probability** P_{sync}: fraction of 50 random initial conditions reaching $r > 0.9$.
3. **Convergence time** t_c: first time $r > 0.9$.

4.2 Results

Experiment 1: Classical benchmark. The measured critical coupling at $\sigma = 1$, $N = 200$ is $K_{\mathrm{c}}^{\mathrm{num}} = 1.55$, compared to the theoretical $K_{\mathrm{c}} = 1.596$ (error $\approx 3\%$, attributable to finite-size effects). At $N = 500$, $K_{\mathrm{c}}^{\mathrm{num}} = 1.632$ (error 2.3%). The square-root scaling $r \sim \sqrt{K_2 - K_{\mathrm{c}}}$ is confirmed. The $\mathcal{O}(N^3) \to \mathcal{O}(N)$ mean-field decomposition of the triplet term enables efficient large-scale simulations.

Experiment 2: Critical slowing down and explosive transition. The relaxation time τ_{relax} near the phase boundary confirms Theorem 1. For $K_3 = 0$ (classical), τ_{relax} peaks at ≈ 12 time units—characteristic of a continuous (second-order) transition with critical slowing down. For $K_3 = 1.5 > K_2^{\mathrm{boundary}}$, τ_{relax} spikes to 48.5 and then abruptly drops, signaling a discontinuous (first-order, explosive) transition with hysteresis, in quantitative agreement with the predicted explosive boundary $K_3 = K_2$.

Experiment 3: Second-order parameter r_2. Our analytical prediction $r_2 = r_1^2$ on the Ott–Antonsen manifold is confirmed quantitatively. At $K_2 = 4$, $K_3 = 2$: $r_2/r_1^2 = 0.972$ (near the OA prediction of 1.0). The ratio systematically deviates for $K_3 < 0$: at $K_2 = 2$, $K_3 = -1$, $r_2/r_1^2 = 0.698$, coinciding with the emergence of 2–3 cluster states. In the incoherent regime ($r_1 \approx 0.1$), $r_2/r_1^2 \gg 1$ (up to ~ 9), confirming that the incoherent state lies far from the OA manifold. These results validate r_2/r_1^2 as a practical diagnostic for OA validity and cluster formation.

Experiment 4: Budget allocation and locking order. Under the constraint $K_2 + K_3 = C$, OA predicts pure pairwise coupling ($K_3 = 0$) maximizes r^*; finite-N simulations show moderate $K_3 > 0$ can improve basin probability instead—confirming that mixed-allocation advantage [Muolo et al., 2025a] is a finite-size effect. Locking order analysis reveals that $K_3 > 0$ weakens the $|\omega_i|$–lock-time correlation (ρ: $-0.55 \to -0.13$), indicating a "winner-takes-all" mechanism where phase alignment overrides frequency proximity.

Experiment 5: Hysteresis and transition order. Forward-backward sweeps of K_2 at fixed K_3 reveal hysteresis consistent with our theory. At $K_3 = 1.5$ (well above the predicted explosive boundary $K_3 \approx K_2^{\mathrm{boundary}}$), the hysteresis gap reaches $\max |r_{\mathrm{fwd}} - r_{\mathrm{bwd}}| = 0.41$—a clear first-order transition. At $K_3 = 1.0$ the gap drops to 0.10, near the predicted boundary. Interestingly, moderate hysteresis ($\sim$ 0.12–0.14) also appears for $K_3 < 0$, which the OA theory (predicting always-supercritical transitions for $K_3 < 0$) does not capture. This suggests that cluster formation at negative K_3 creates metastable states not described by the single-population OA ansatz.

Experiments 6–7: Chimera search and distribution effects. No chimera states were found across 15 (K_2, K_3) configurations ($N = 500$), confirming that higher-order interactions alone cannot generate chimeras in all-to-all networks. Comparing Gaussian, Lorentzian, and uniform distributions: K_c depends on $g(0)$ as predicted (Gaussian: 1.28, uniform: 1.81), and K_3-sensitivity is weak ($< 10\%$), consistent with Proposition 1.

Experiment 8: 2D phase diagram. The 20×20 scan over $K_2 \in [0, 4]$, $K_3 \in [-2, 2]$ at $\sigma = 1$, $N = 200$ provides the first complete numerical phase diagram for this model. Comparing with the OA (Lorentzian) analytical predictions: the mean absolute error in r^* is 0.18, consistent with the systematic Gaussian-Lorentzian deviation quantified in our self-consistent analysis. The explosive synchronization boundary is detected at $(K_2, K_3) \approx (1.47, 1.58)$, close to the predicted diagonal $K_3 = K_2$ (Theorem 1). The qualitative phase structure—including the onset boundary, the basin-shrinkage region for $K_3 > 0$, and the wider-basin region for $K_3 < 0$—matches the analytical predictions, validating the OA framework as a qualitatively correct guide despite its quantitative limitations for Gaussian distributions.

Experiment 9: Re-entrant region. Our OA theory predicts re-entrant synchronization for K_2 slightly above K_c when K_3 exceeds the explosive boundary $K_3 \approx K_2$. In the 20×20 scan with $K_3 \in [-2, 2]$, r^* increases monotonically with K_3—the re-entrant drop is not observed because $K_3^{\max} = 2$ does not extend far enough into the explosive regime. However, the hysteresis experiment (Exp. 5) at $K_3 = 1.5$ shows hallmarks of first-order transition (gap 0.41), and the critical slowing data (Exp. 2) at $K_3 = 1.5$ shows an abrupt relaxation time spike followed by collapse—both signatures consistent with the onset of subcritical dynamics that would produce re-entrance at larger K_3/K_2 ratios. Targeted verification with $K_3 \in [0, 6]$ at $K_2 \approx K_c$ is left to future work.

5 Discussion

Our unified framework resolves the apparent contradictions in the literature through a simple geometric insight: the higher-order coupling K_3 reshapes the entire *landscape* of the radial potential

$$V(r) = -\int_0^r \dot{r}' \, \mathrm{d}r' = \frac{\Delta}{2} r^2 - \frac{K_2}{2}\left(\frac{r^2}{2} - \frac{r^4}{4}\right) - \frac{K_3}{2}\left(\frac{r^4}{4} - \frac{r^6}{6}\right), \tag{23}$$

and different studies were measuring different features of this landscape:

- **Onset** (K_c): Curvature at $r = 0 \to$ independent of K_3 (Proposition 1).
- **Basin**: Position of the unstable saddle $\to K_3$ raises it (Proposition 2).
- **Stability**: Curvature at $r^* \to K_3$ deepens the well for moderate values.
- **Transition type**: Sign of cubic coefficient $\to K_3 > K_2$ makes it explosive (Theorem 1).

The re-entrant region (Section 3.7) is a genuinely new prediction that has not been reported in the literature and can be tested experimentally in coupled electrochemical oscillators or Josephson junction arrays.

5.1 Asymmetry between positive and negative K_3

Our analysis reveals a fundamental asymmetry between cooperative ($K_3 > 0$) and competitive ($K_3 < 0$) higher-order interactions:

- $K_3 > 0$: "Elite synchronization"—stronger phase locking ($r^* \uparrow$) but harder to achieve (basin$\downarrow$), with potentially explosive transitions.
- $K_3 < 0$: "Democratic synchronization"—weaker but more robust synchrony ($r^* \downarrow$, basin$\uparrow$), always continuous transitions on the OA manifold.

This asymmetry suggests that biological systems requiring robust, moderate-amplitude synchrony (e.g., cortical oscillations) may preferentially employ competitive higher-order coupling, while systems requiring strong phase locking (e.g., cardiac pacemakers) benefit from cooperative coupling. Numerically, the $K_3 < 0$ regime additionally supports multi-cluster states—a beyond-OA phenomenon requiring multi-population extensions.

5.2 Connection to time-delayed systems

Our framework also illuminates the recently discovered equivalence between time-delayed and higher-order coupling [Muolo et al., 2025b]. A Taylor expansion of the delayed Kuramoto model to $\mathcal{O}(\tau)$ yields effective parameters (see Appendix B for the full derivation):

$$K_2^{\text{eff}} = K_2 - \frac{K_2^2\tau}{2}, \qquad K_3^{\text{eff}} = \frac{K_2^2\tau}{4}. \tag{24}$$

Three important consequences follow:

1. $K_3^{\text{eff}} \geq 0$ always: time delay can only induce *positive* three-body coupling. The $K_3 < 0$ half of our phase diagram is inaccessible to purely delayed systems.
2. The physically reachable region in the $(K_2^{\text{eff}}, K_3^{\text{eff}})$ plane is a triangle bounded by $K_3 = 0$ and $K_3 = (K_2 - K_2^{\text{eff}})/2$.
3. The explosive synchronization line $K_3 = K_2$ (Theorem 1) lies outside this triangle, explaining why explosive transitions are *not* typically observed in time-delayed Kuramoto systems.

5.3 Limitations

Our three-body truncation is justified when $K_4 r^4/(K_3 r^2 + K_2) \ll 1$, which holds for $K_3 < K_2/2$ and $r^* < 0.7$ (error $< 5\%$); near the explosive boundary, four-body terms ($\sim$12% effect at $N = 50$) should be included. The OA reduction also assumes all-to-all coupling; sparse network extensions [Muolo et al., 2025a] remain an important direction.

6 Related Work

Classical Kuramoto theory. The original model was introduced by Kuramoto [1984]. The Ott–Antonsen reduction [Ott and Antonsen, 2008] enabled exact low-dimensional analysis.

Higher-order interactions. Skardal and Arenas [2020] first demonstrated explosive synchronization with higher-order coupling. Zhang et al. [2024] showed basin shrinkage via quasipotential analysis. Muolo et al. [2025a] studied random hypergraphs and found a non-monotonic dependence. Wang et al. [2025] analyzed ring networks with moderate coupling. Dai and Kori [2025] found up to 9 synchronization transitions in three-oscillator systems. The connection between time-delayed and higher-order coupling was established by Muolo et al. [2025b].

7 Conclusion

We have provided a unified analytical framework for understanding the effects of higher-order interactions on Kuramoto synchronization. The key contributions are:

1. The first complete Ott–Antonsen reduction for the three-body Kuramoto model, yielding a scalar ODE (14) for the order parameter.
2. Three analytically derived critical manifolds that reconcile the four major contradictory studies (2020–2025).
3. Discovery of a re-entrant synchronization region.
4. A σ-scaling law (Proposition 3) showing phase boundaries are self-similar, answering the open question of Muolo et al. [2025a].
5. Explicit mapping $K_3^{\text{eff}} = K_2^2\tau/4$ from time-delayed to higher-order models, delineating the physically reachable region.
6. Quantitative validation via $> 10^5$ GPU-accelerated simulations.

References

Qionglin Dai and Hiroshi Kori. Higher-order synchronization and multistability in higher-order Kuramoto phase oscillators. *Nonlinear Dynamics*, 113:21801–21811, 2025.

Yoshiki Kuramoto. *Chemical Oscillations, Waves and Turbulence*. Springer, 1984. doi: 10.1007/978-3-642-69689-3.

Riccardo Muolo, Hiroya Nakao, and Marco Coraggio. When higher-order interactions enhance synchronization: the case of the Kuramoto model on random hypergraphs. *arXiv preprint*, 2025a.

Riccardo Muolo et al. Emergence of higher-order interactions in systems of coupled Kuramoto oscillators with time delay. *arXiv preprint*, 2025b.

Edward Ott and Thomas M. Antonsen. Low dimensional behavior of large systems of globally coupled oscillators. *Chaos*, 18:037113, 2008. doi: 10.1063/1.2930766.

Per Sebastian Skardal and Alex Arenas. Higher order interactions in complex networks of phase oscillators promote abrupt synchronization switching. *Communications Physics*, 3:218, 2020. doi: 10.1038/s42005-020-00485-0.

Zijian Wang, Wentao Qi, Jianfeng Zhu, and Xin Liu. Moderate higher-order interactions enhance stability while preserving basin structure. *arXiv preprint*, 2025.

Yuanzhao Zhang, Per Sebastian Skardal, Federico Battiston, Giovanni Petri, and Maxime Lucas. Deeper but smaller: Higher-order interactions increase linear stability but shrink basins. *Science Advances*, 10:eado8049, 2024. doi: 10.1126/sciadv.ado8049.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer blue[Yes], orange[No], or gray[N/A].
- gray[N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for gray[N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While blue[Yes] is generally preferable to orange[No], it is perfectly acceptable to answer orange[No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering orange[No] or gray[N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer blue[Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading "NeurIPS Paper Checklist"**,
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers**.

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: red**[TODO]**

   Justification: red**[TODO]**

   Guidelines:

   - The answer gray[N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A orange[No] or gray[N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: red**[TODO]**

Justification: red**[TODO]**

Guidelines:

- The answer gray[N/A] means that the paper has no limitation while the answer orange[No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: red**[TODO]**

   Justification: red**[TODO]**

   Guidelines:

   - The answer gray[N/A] means that the paper does not include theoretical results.
   - All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
   - All assumptions should be clearly stated or referenced in the statement of any theorems.
   - The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
   - Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
   - Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

   Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

   Answer: red**[TODO]**

Justification: red**[TODO]**

Guidelines:

- The answer gray[N/A] means that the paper does not include experiments.
- If the paper includes experiments, a orange[No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

   Answer: red**[TODO]**

   Justification: red**[TODO]**

   Guidelines:

   - The answer gray[N/A] means that paper does not include experiments requiring code.
   - Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
   - While we encourage the release of code and data, we understand that this might not be possible, so orange[No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
   - The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
   - The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
   - The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.

- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: red**[TODO]**

   Justification: red**[TODO]**

   Guidelines:

   - The answer gray[N/A] means that the paper does not include experiments.
   - The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
   - The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: red**[TODO]**

   Justification: red**[TODO]**

   Guidelines:

   - The answer gray[N/A] means that the paper does not include experiments.
   - The authors should answer blue[Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
   - The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
   - The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
   - The assumptions made should be given (e.g., Normally distributed errors).
   - It should be clear whether the error bar is the standard deviation or the standard error of the mean.
   - It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
   - For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
   - If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: red**[TODO]**

   Justification: red**[TODO]**

   Guidelines:

   - The answer gray[N/A] means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.

- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

   Answer: red**[TODO]**

   Justification: red**[TODO]**

   Guidelines:

   - The answer gray[N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
   - If the authors answer orange[No], they should explain the special circumstances that require a deviation from the Code of Ethics.
   - The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: red**[TODO]**

    Justification: red**[TODO]**

    Guidelines:

    - The answer gray[N/A] means that there is no societal impact of the work performed.
    - If the authors answer gray[N/A] or orange[No], they should explain why their work has no societal impact or why the paper does not address societal impact.
    - Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
    - The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
    - The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
    - If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

    Answer: red**[TODO]**

    Justification: red**[TODO]**

    Guidelines:

- The answer gray[N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: red**[TODO]**

    Justification: red**[TODO]**

    Guidelines:

    - The answer gray[N/A] means that the paper does not use existing assets.
    - The authors should cite the original paper that produced the code package or dataset.
    - The authors should state which version of the asset is used and, if possible, include a URL.
    - The name of the license (e.g., CC-BY 4.0) should be included for each asset.
    - For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
    - If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
    - For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
    - If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: red**[TODO]**

    Justification: red**[TODO]**

    Guidelines:

    - The answer gray[N/A] means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: red**[TODO]**

Justification: red**[TODO]**

Guidelines:

- The answer gray[N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: red**[TODO]**

    Justification: red**[TODO]**

    Guidelines:

    - The answer gray[N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
    - Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
    - We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
    - For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

    Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

    Answer: red**[TODO]**

    Justification: red**[TODO]**

    Guidelines:

    - The answer gray[N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
    - Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

A Detailed Derivations

A.1 Mean-field reduction of the three-body term

Starting from the discrete three-body coupling:

$$T_i = \frac{K_3}{N^2} \sum_{j=1}^{N} \sum_{k=1}^{N} \sin(2\theta_j - \theta_k - \theta_i). \tag{25}$$

Using the identity $\sin(A) = \mathrm{Im}[\mathrm{e}^{\mathrm{i}A}]$:

$$T_i = \frac{K_3}{N^2}\,\mathrm{Im}\left[\mathrm{e}^{-\mathrm{i}\theta_i}\sum_{j=1}^{N}\mathrm{e}^{2\mathrm{i}\theta_j}\sum_{k=1}^{N}\mathrm{e}^{-\mathrm{i}\theta_k}\right]$$
$$= K_3\,\mathrm{Im}\big[\mathrm{e}^{-\mathrm{i}\theta_i}\,Z_2\,\bar{Z}_1\big], \tag{26}$$

where $Z_n = \frac{1}{N}\sum_j \mathrm{e}^{\mathrm{i}n\theta_j}$.

In the $N \to \infty$ limit, $Z_n \to \int \mathrm{e}^{\mathrm{i}n\theta} f(\theta,\omega,t)\,\mathrm{d}\theta\,\mathrm{d}\omega$, confirming (7).

A.2 Derivation of the α-ODE

Substituting the Ott–Antonsen ansatz $\hat{f}_n = \alpha^n$ into the continuity equation and using the velocity field (10):

The n-th Fourier component of the continuity equation gives

$$\dot{\hat{f}}_n = -\mathrm{i}n\omega\,\hat{f}_n + \frac{n}{2}\big[H\,\hat{f}_{n-1} - \bar{H}\,\hat{f}_{n+1}\big], \tag{27}$$

where $H = K_2 Z_1 + K_3 Z_2 \bar{Z}_1$ is the effective coupling field.

On the Ott–Antonsen manifold, $\hat{f}_n = \alpha^n$, so

$$n\alpha^{n-1}\dot{\alpha} = -\mathrm{i}n\omega\,\alpha^n + \frac{n}{2}\big[H\,\alpha^{n-1} - \bar{H}\,\alpha^{n+1}\big]. \tag{28}$$

Dividing by $n\alpha^{n-1}$ yields (11).

A.3 Radial dynamics derivation

Setting $\omega_0 = 0$, evaluating α at $\omega = -\mathrm{i}\Delta$ for the Lorentzian, and writing $z = \bar{\alpha}(-\mathrm{i}\Delta)$ with $z = r\mathrm{e}^{\mathrm{i}\phi}$:

From (11):

$$\dot{\bar{\alpha}} = \mathrm{i}\omega\bar{\alpha} + \frac{1}{2}\big[\bar{H} - H\,\bar{\alpha}^2\big]. \tag{29}$$

At $\omega = -\mathrm{i}\Delta$, $\bar{\alpha} = z$, $Z_n = z^n$:

$$\dot{z} = -\Delta z + \frac{1}{2}\big[\bar{H} - Hz^2\big] \tag{30}$$
$$= -\Delta z + \frac{1}{2}\big[(K_2\bar{z} + K_3\bar{z}^2 z) - (K_2 z + K_3 z^2\bar{z})z^2\big]. \tag{31}$$

With $z = r\mathrm{e}^{\mathrm{i}\phi}$ and extracting the radial part:

$$\dot{r} = -\Delta r + \frac{1}{2}\big[K_2 r(1-r^2) + K_3 r^3(1-r^2)\big] \tag{32}$$
$$= -\Delta r + \frac{r(1-r^2)}{2}(K_2 + K_3 r^2), \tag{33}$$

confirming (14). □

B Time-Delay Equivalence

The full derivation of the time-delay to higher-order mapping is provided in the supplementary file `time_delay_equivalence.tex`. Here we state the key result.

Starting from the delayed Kuramoto model $\dot{\theta}_i = \omega_i + \frac{K_2}{N}\sum_j \sin(\theta_j(t-\tau) - \theta_i(t))$, a Taylor expansion to $\mathcal{O}(\tau)$ followed by trigonometric product-to-sum identities yields:

$$K_2^{\text{eff}} = K_2 - \frac{K_2^2\tau}{2}, \tag{34}$$

$$K_3^{\text{eff}} = \frac{K_2^2\tau}{4}. \tag{35}$$

The physically reachable region satisfies $K_3^{\text{eff}} = (K_2 - K_2^{\text{eff}})/2$ with $K_3^{\text{eff}} \geq 0$, forming a triangle in the $(K_2^{\text{eff}}, K_3^{\text{eff}})$ plane.